

Deep learning

Neural networks

Matthew Greenberg

May 2022

This course

1. Introduction
2. Convolutional neural networks (CNNs) for computer vision
3. Text processing and recurrent neural networks (RNNs)

Deep learning

- ▶ Probabilistic models based on *neural networks*.
- ▶ Transformative for *computer vision*, *natural language processing*, automation, robotics, molecular modeling, simulation...
- ▶ Benefits of deep learning most pronounced for analysis of *unstructured, high-dimensional data*.

Using deep learning powered APIs and packages


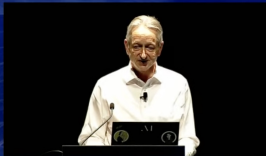
- ▶ Many companies sell access deep learning powered APIs for common data processing and analysis tasks.
- ▶ Let's look at how to use Google's cloud vision API...
- ▶ Instead of accessing the model via API, you can download it and incorporate it into your code.
- ▶ Let's download a transformer model for extractive question answering from HuggingFace...

Artificial intelligence (AI)

Geoffrey Hinton and Yann LeCun, 2018 ACM A.M. Turing Award Lecture "The Deep Learning Revolution"

Two ways to make a computer do what you want

- **Intelligent design:** Figure out consciously exactly how you would manipulate symbolic representations to perform the task and then tell the computer, in excruciating detail, exactly what to do.
- **Learning:** Show the computer lots of examples of inputs together with the desired outputs. Let the computer learn how to map inputs to outputs using a general purpose, learning procedure.



<https://www.youtube.com/watch?v=VsnQf7exv5I> (11:49)

Machine learning (ML)

Show the computer lots of examples of inputs together with the desired outputs. Let the computer learn how to map inputs to outputs using a general purpose learning procedure.

input

output

image

label



deer



frog



truck

input

output

English text

French translation

Can I borrow your magazine?

Puis-je emprunter votre magazine?

Please pass the ketchup.

Veuillez passer le ketchup.

His shirt was made in turkey.

Sa chemise a été confectionnée en Turquie.

Deep learning (DL)



Yann LeCun @ylecun · Dec 24, 2019



Some folks still seem confused about what deep learning is. Here is a definition:

DL is constructing networks of parameterized functional modules & training them from examples using gradient-based optimization....

facebook.com/722677142/post...

💬 37


↻ 551

❤️ 1.8K

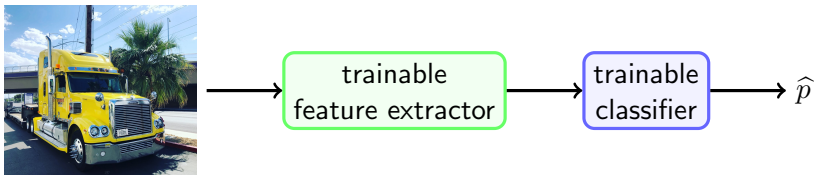


Networks

Traditional machine learning

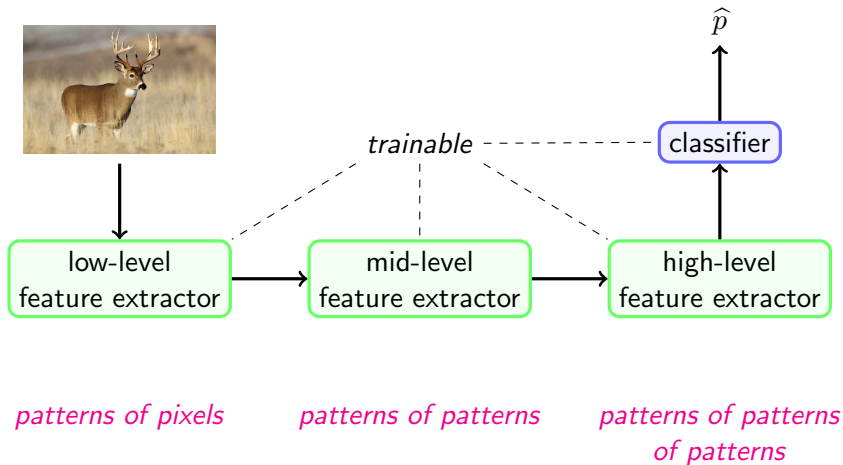


Deep learning



- \hat{p} is a vector of class probabilities.

Feature hierarchy




Linear layers


Trainable parameters:

► **weights:** $W \in \mathbb{R}^{p \times q}$

► **biases:** $b \in \mathbb{R}^q$




More concisely:




$$v = L_{W,b}(u) = uW + b$$

Compositions of linear layers are linear





$$\begin{aligned} L_{W_2, b_2}(L_{W_1, b_1}(u)) &= L_{W_2, b_2}(uW_1 + b_1) \\ &= (uW_1 + b_1)W_2 + b_2 \\ &= u \underbrace{W_1 W_2}_W + \underbrace{b_1 W_2 + b_2}_b \\ &= L_{W, b}(u) \end{aligned}$$



ReLU-activated linear layers

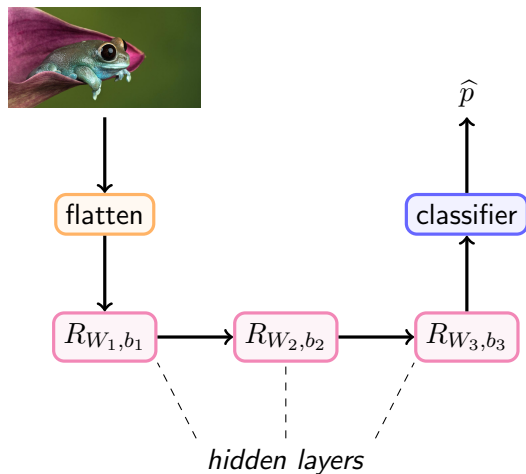
Rectified linear unit:

$$\text{ReLU}(x) = \max(x, 0)$$




$$R_{W,b}(u) = \text{ReLU} (L_{W,b}(u)) \quad (\text{apply ReLU componentwise})$$

A simple classification network




- The **flatten** layer flattens an RGB image tensor with shape $m \times n \times c$ to an mnc -dimensional vector.

The classification layer: softmax-activated linear layer



- The classification layer of a K -class classification network typically performs **logistic regression**, mapping features output by the hidden layers onto class probabilities.



The softmax function

- ▶ The **softmax function** converts “raw scores” u_0, \dots, u_{p-1} into a **distribution vector**, i.e. a vector with nonnegative components that sum to 1:

$$\text{softmax}(u_0, \dots, u_{p-1}) = \frac{1}{\sum_j e^{u_j}} (e^{u_0}, \dots, e^{u_{p-1}})$$

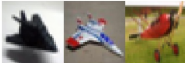

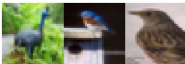
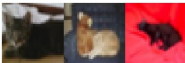

Using pretrained image classification models






- ▶ You can download many pretrained image classification models.
- ▶ Most are trained on the **ImageNet dataset**, with
 - ▶ 1000 object classes,
 - ▶ > 1.2 million training images




Example: the CIFAR-10 dataset

- ▶ 60000 RGB images, 32×32 pixels
- ▶ 10 classes

label/class	examples
0/airplane	
1/automobile	
2/bird	
3/cat	
4/deer	

label/class	examples
5/dog	
6/frog	
7/horse	
8/ship	
9/truck	

An image classifier for CIFAR-10



- ▶ input x : tensor with shape $32 \times 32 \times 3$
- ▶ target: $y \in \{0, 1, \dots, 9\}$
- ▶ output \hat{p} : vector of length 10, where


$$\hat{p}_j = \text{probability that } x \text{ has class } j$$

- ▶ trainable parameters: θ , initialized randomly

Training the classifier

- ▶ loss function: penalize misclassifications

$$L(y, \hat{p}) = -\log \hat{p}_y$$



- ▶ Adjust parameters to decrease loss via

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} L(y, \hat{p})$$

(α is a positive hyperparameter called the **learning rate**.)

Batches and epochs

- ▶ One training **epoch**:
 - ▶ Partition training data into N/I **batches** of size I .
 - ▶ for $0 \leq k \leq N/I$:

- ▶ Compute

$$\hat{p}_i = f_{\theta}(x_i), \quad kI \leq i < (k+1)I.$$

- ▶ Adjust parameters to decrease average loss over the batch via

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} \left\{ \frac{1}{I} \sum_{bI \leq i < (b+1)I} L(y_i, \hat{p}_i) \right\}.$$

- ▶ Train multiple epochs, until convergence.

```
model.fit(x_train, y_train, batch_size=100, epochs=20);
```

Epoch 1/20

500/500 [=====] - 1s 2ms/step - loss: 2.0992 - accuracy: 0.2314

Epoch 2/20

500/500 [=====] - 1s 2ms/step - loss: 1.8704 - accuracy: 0.3424

Epoch 3/20

500/500 [=====] - 1s 2ms/step - loss: 1.8263 - accuracy: 0.3621

⋮

Epoch 18/20

500/500 [=====] - 1s 2ms/step - loss: 1.7248 - accuracy: 0.4058

Epoch 19/20

500/500 [=====] - 1s 2ms/step - loss: 1.7164 - accuracy: 0.4129

Epoch 20/20

500/500 [=====] - 1s 2ms/step - loss: 1.7209 - accuracy: 0.4062

